Multi-Step Wizard Design for Esteemed Ventures Cost Estimator

Based on the provided business and service details from Esteemed Ventures' documents (e.g., the 22+ specialized autonomous AI agents, investment lifecycle workflows, and operational structure), I've designed a **multi-step wizard** for feeding customer data into the cost estimator. This wizard guides users through inputting task-specific details while auto-populating defaults from **predefined task presets**. The presets are derived directly from your AI-native platform's key domains: Investment Ops, Fund Admin & IR, Marketing & Origination, Deal Execution, Post-Close Value Creation, and Back Office.

The wizard ensures comprehensiveness by mapping to the core data categories (Task Details, Agent Configuration, Usage Metrics, Cost Structure, External Factors) while keeping it user-friendly—e.g., via dropdowns, sliders, and conditional logic. It outputs a detailed cost breakdown (e.g., per-agent token usage, total OpEx savings vs. traditional models) and supports "what-if" scenarios.

Wizard Overview
- **Entry Point**: User selects a **preset** (or starts custom) based on common workflows like "Deal Sourcing" or "Portfolio Optimization."
- **Steps**: 5 progressive screens, with progress bar and "Back" navigation. Defaults load from preset; users can override.
- **Validation**: Real-time checks (e.g., agent compatibility with task type) and tooltips linking to agent descriptions.
- **Output**: Interactive summary with pie chart (agent costs), sensitivity sliders (e.g., adjust error rate), and export (PDF/CSV). Includes 60-70% OpEx savings benchmark vs. traditional PE.
- **Tech Notes**: Implement in React/Vue for frontend; backend pulls from your SLM/Enterprise Memory System for dynamic defaults. Integrate MCP for real-time pricing from LLMs (e.g., OpenAI/Claude/Groq rates).

Step-by-Step Wizard Flow

Step	Screen Title	Key Inputs Collected	Defaults from Presets	Conditional Logic/UI Tips
1: Task Selection	"Define Your Workflow"	- Task Type/Category (dropdown: Investment Ops, Fund Admin & IR, etc.) - Task Complexity (slider: Low/Med/High or 1-10) - Expected Duration/Volume (e.g., hours, # of deals/subtasks)	- Preset auto-selects category and complexity (e.g., "Deal Sourcing" defaults to Med complexity, 4-8 hours). - Volume: Scaled to domain (e.g., 5 deals for sourcing).	- Preset dropdown with 8 options (see below). - If custom, show agent compatibility warning. - Tooltip: "Based on Esteemed Agents' modular design for parallel processing."
2: Agent Configuration	"Assemble Your Agent Team"	- Agent

Types (multi-select: e.g., Clay for PM, Finley for Modeling; up to 5-10 based on task).
- # of Agents (auto-suggest min/max).
- Capabilities/Limits (e.g., tokens per call, error rate %). | - Preset loads 3-6 agents (e.g., "Financial Modeling" defaults: Finley + Clay + Lex).
- Limits: Defaults to agent specs (e.g., FinGPT for finance tasks, 10K tokens/call). | - Searchable agent list with descriptions (e.g., "Finley: Harvard-level modeling").
- Auto-add backups (e.g., Ivy for compliance).
- Slider for parallelism (1-5x speed vs. cost multiplier). |
3: Usage & Performance	"Estimate Resource Needs"	- Resource Consumption (e.g., API calls, data transfer MB; auto-calc from duration). - Efficiency Factors (e.g., error rate 5%, retry multiplier). - Scalability (e.g., batch size, auto-scale threshold).	- Preset: Efficiency at 90% (based on continual learning); e.g., "Legal Review" defaults 20 API calls, 5% error. - Batch: On for high-volume (e.g., 10+ subtasks).	- Auto-fill from Step 1 volume (e.g., High complexity adds 20% to calls). - Progress bar for historical benchmarks (e.g., "Similar tasks avg. 85% success"). - Checkbox: "Use multi-model routing for cost optimization."
4: Cost & Financials	"Set Pricing & Billing"	- Per-Agent Pricing (e.g., $0.01/1K tokens; tiered). - Billing Model (pay-per-use, subscription, hybrid). - Taxes/Fees/Margins (e.g., 20% markup, region VAT).	- Preset: Token + seat hybrid (e.g., $29/month/seat + tokens; 85% gross margin target). - Discounts: Auto-apply for volume (e.g., 10% off for >10 agents).	- Dropdown for tiers (Individual/Dept/Enterprise). - Real-time total preview. - Toggle: "BYO API keys" for 30% markup reduction.
5: Review & Risks	"Finalize & Simulate"	- External Factors (e.g., inflation 3%, risk buffer 15%). - User Profile (e.g., Enterprise vs. HNWI). - Success Metrics (e.g., >95% accuracy).	- Preset: Buffer at 10% for deal velocity; profile defaults to "Institutional LP." - Metrics: Tied to task (e.g., 3-5x ROI projection for PE fund sim).	- Sensitivity sliders (e.g., "+1 agent = +$X"). - Risk summary table (e.g., "Tech Risk: Mitigated by multi-model"). - Button: "Run Simulation" → Generates report with 12-18% IRR projection for credit tasks.

Predefined Task Presets
These 8 presets cover 80%+ of your investment lifecycle (from sourcing to value creation), based on the 22 agents and workflows in the docs. Each preset pre-fills Steps 1-4 with realistic defaults, drawing from your 60-70% OpEx savings model, multi-model architecture (OpenAI/Claude/Groq), and continual learning cadence. Presets can be expanded via your SLM for custom firm-specific playbooks.

Preset Name	Domain	Description	Default Complexity/Duration	Default Agents (3-6)	Default Usage (e.g., Calls/Tokens)	Default Cost Model	Projected Savings vs. Traditional

------------------|-----------------------------------|-------------
-------|---------------------------------|
| **Deal Sourcing** | Investment Ops | Autonomous sourcing & pipeline mgmt for roll-ups/SaaS targets. | Med / 4–8 hrs | Clay (PM), Lyla (Outreach), Gemma (CTO) | 15 calls, 50K tokens; 5% error | Token + seat ($49/month + $0.01/1K) | 65% (vs. $150K assoc. salary) |
| **Financial Modeling** | Investment Ops | Qualify deals, build models for EBITDA analysis. | High / 6–12 hrs | Finley (Modeling), Clay (PM), FinanceBot | 25 calls, 100K tokens; batch=10 | Hybrid ($149/dept + tokens) | 70% (Harvard-level output, no CapEx) |
| **Legal Review** | Deal Execution | Dataroom prep, term sheets, redlines. | Med / 3–6 hrs | Lex (Legal), Ivy (Compliance), Bree (IR) | 10 calls, 30K tokens; real-time routing | Pay-per-use ($0.005/1K via Groq) | 60% (parallel vs. sequential human) |
| **Investor Relations** | Fund Admin & IR | LP comms, reporting, subscription docs. | Low / 2–4 hrs | Bree (IR), Pax (FundOps), Mia (MarTech) | 8 calls, 20K tokens; weekly retrain | Subscription ($29/seat) | 75% (automated portals vs. $200–300K headcount) |
| **GTM Playbook Implementation** | Post-Close Value Creation | Post-acq. GTM, integrations for staffing/SaaS. | High / 8–16 hrs | Ava (AI Ops), Jax (Product), Kai (DevOps) | 30 calls, 150K tokens; Vision AI add-on | Enterprise ($999 + unlimited) | 68% (synergies via pgvector memory) |
| **Portfolio Optimization** | Post-Close Value Creation | Real-time monitoring, synergy ID across holdings. | Med / 4–10 hrs | Ava (AI Ops), Finley (Modeling), Chibs (DevOps) | 20 calls, 80K tokens; daily ingestion | Token pool ($299/dept) | 70% (continuous vs. quarterly reviews) |
| **Marketing Campaign** | Marketing & Origination | LP/founder outreach, event follow-up. | Low / 2–5 hrs | Lyla (Outreach), Mia (MarTech), Happy (QA) | 12 calls, 40K tokens; HeyGen integration | Pay-per-use + 10% discount | 62% (personalized avatars vs. SDR team) |
| **Back Office Ops** | Back Office | Infrastructure monitoring, financial ops, QA. | Low / 1–3 hrs | Chibs (DevOps), Happy (QA), FinanceBot | 5 calls, 15K tokens; monthly fine-tune | Flat ($99/month) | 80% (24/7 vs. FTE per co.) |

Implementation Recommendations
- **Customization**: Allow users to "Fork" a preset and add agents (e.g., spin up task-specific swarms for scale).
- **Dynamic Elements**: Pull live LLM rates via API (e.g., Groq for speed). Use your SLM for preset evolution (e.g., weekly retrain on usage logs).
- **Edge Cases**: Handle high-scale (e.g., 100+ deals) with auto-scaling warnings; include risk buffers for market volatility (e.g., 3% inflation).
- **Metrics Tracking**: Post-wizard, log anonymized data to your Enterprise Memory System for LTV:CAC optimization (target 56:1).
- **Testing**: Prototype with sample runs (e.g., "Deal Sourcing" preset yields ~$50–150 total, 85% margin).

This setup makes the estimator comprehensive yet intuitive, aligning with your AI-native ethos. If you provide sample pricing data or want code snippets for the wizard logic, I can refine further!

TASK: Esteemed Ventures (EV) SME Agents Service

Build an multi-agent swarm estimator that helps users know what the project may cost, and time expected for completion using Esteemed Ventures (EV) SME Agents Service

EV Core Agent Swarm Cost Estimator should be based on:

1. Task
2. Scale
3. Performance

Other items: task characteristics, agent specifics, usage patterns, operational overheads, and financial variables. This ensures the app can model real-world scenarios accurately, handle variability, and provide actionable insights (e.g., cost breakdowns, optimization suggestions).

Core Information to Gather for a Comprehensive Cost Calculator App

To build a robust cost calculator for a multi-agent service, you'll need to collect data across several dimensions: task characteristics, agent specifics, usage patterns, operational overheads, and financial variables. This ensures the app can model real-world scenarios accurately, handle variability, and provide actionable insights (e.g., cost breakdowns, optimization suggestions).

I've organized the key information into categories below, with examples of why it's important and how it might be input into the app. Aim to make inputs user-friendly (e.g., dropdowns for agent types, sliders for task complexity) and allow for custom overrides.

1. **Task Details**
   These define the scope and drive how agents are allocated.
   - **Task Type/Category**: E.g., data analysis, content generation, customer support. *Why?* Different tasks require specific agent combinations, affecting cost baselines.
   - **Task Complexity Level**: Low/medium/high or a numerical scale (1-10). *Why?* Higher complexity may need more agents or iterations.
   - **Expected Duration or Volume**: E.g., hours/days, number of items processed (like emails or queries). *Why?* Scales resource usage linearly or exponentially.
   - **Dependencies/Sub-tasks**: Breakdown of subtasks and their sequence. *Why?* Reveals handoffs between agents, which incur coordination costs.

   - **Success Metrics/Quality Threshold**: E.g., accuracy rate >95%. *Why?* Low success may trigger retries, inflating costs.

2. **Agent Configuration**
   Focus on the "workers" in your system.
   - **Agent Types Available**: List of agents (e.g., LLM for reasoning, API caller for integrations, validator for quality checks) with descriptions. *Why?* Users select based on task fit; each has unique pricing.
   - **Number of Agents per Task**: Minimum/maximum, or dynamic based on load. *Why?* Parallel agents speed up tasks but multiply costs.
   - **Agent Capabilities and Limits**: E.g., tokens per call for AI agents, API rate limits. *Why?* Prevents overestimation; helps simulate bottlenecks.
   - **Agent Assignment Rules**: E.g., primary agent + backups. *Why?* Models failover scenarios for reliability.

3. **Usage and Performance Metrics**
   Track runtime consumption for precise calculations.
   - **Resource Consumption Rates**: E.g., compute units (GPU hours), API calls, data transfer (MB). *Why?* Core drivers of variable costs in cloud-based agents.
   - **Efficiency Factors**: E.g., average time per subtask, error rate (5–10%). *Why?* Accounts for waste; enables "what-if" simulations (e.g., "What if error rate drops?").
   - **Scalability Options**: E.g., auto-scaling thresholds, batch processing discounts. *Why?* For high-volume tasks, costs per unit decrease.
   - **Historical Data**: Past task logs (if available) for benchmarking. *Why?* Improves accuracy over time via machine learning refinements.

4. **Cost Structure and Financial Inputs**
   Define how costs are computed and billed.
   - **Per-Agent Pricing**: E.g., $0.01 per 1K tokens, $5/hour flat rate. *Why?* Foundation of all calculations; include tiers (e.g., volume discounts).
   - **Bundled or Shared Costs**: E.g., platform fees, shared infrastructure (10% overhead). *Why?* Captures non-agent expenses like orchestration tools.
   - **Billing Model**: Pay-per-use, subscription, or hybrid. *Why?* Affects total cost formulas (e.g., amortized fixed costs).
   - **Taxes, Fees, and Margins**: E.g., VAT rates by region, service markup (20%). *Why?* For end-to-end pricing; make region-specific.
   - **Discounts/Promotions**: E.g., first-task free, loyalty tiers. *Why?* Encourages user adoption; dynamic based on volume.

5. **External and Risk Factors**
   Add layers for realism and forecasting.
   - **Market/Inflation Adjustments**: E.g., annual cost increases

(3–5%). *Why?* Future-proofs estimates for long-term planning.
   - **Risk Buffers**: E.g., contingency for delays (add 15% buffer). *Why?* Handles uncertainties like agent downtime.
   - **User Profile**: E.g., enterprise vs. individual, expected usage frequency. *Why?* Tailors recommendations (e.g., bulk pricing for enterprises).
   - **Integration Hooks**: E.g., API keys for real-time pricing feeds. *Why?* Keeps data fresh without manual updates.

Implementation Tips for Your App
- **Data Collection Method**: Use a multi-step wizard for inputs, with defaults based on task presets (e.g., "Standard email campaign" auto-fills agents and rates).
- **Output Features**: Beyond totals, show breakdowns (pie charts for agent costs), sensitivity analysis (e.g., "Increasing agents by 1 adds $X"), and exportable reports.
- **Validation and Iteration**: Start with MVP data (e.g., just task + agents), then expand based on user feedback. Test with sample scenarios to ensure edge cases (e.g., zero-cost tasks) are handled.
- **Privacy Note**: If gathering historical data, ensure compliance (e.g., anonymize usage logs).

This framework should cover 90%+ of scenarios—start gathering this info via user surveys or prototypes to refine it further! If you share more details about the agents (e.g., AI vs. human), I can tailor it more.